



# MLOps and LLM Observability


Karthikeyan Vaiyapuri



01

# Overview

# ● Topics




- 01 Evaluating model performance and metrics
- 02 Monitoring during training and inference
- 03 Model deployment basics with Flask/Streamlit
- 04 Hands-on lab: Deploy a simple model



# 02


## Introduction to Model Validation Techniques

# ● Introduction to Model Validation Techniques




## What is model validation?

- Assessing how well your model generalizes to unseen data
- Prevents overfitting and underfitting



## Common Techniques:

- **Hold-out validation:** Split data into train/test sets (e.g., 80/20)
- **K-fold cross-validation:** Rotate through K train/test splits for robust estimates
- **Stratified sampling:** Maintains class proportions in splits



**Goal:** Ensure your model performs well on real, unseen data.





03

## Evaluating Model Performance and Metrics

# ● Evaluating Model Performance and Metrics

## Classification

(Metrics)

- **Accuracy:** Overall correctness of predictions
- **Precision/Recall/F1:** Quality and coverage of positive predictions
- **AUC-ROC:** Probability curve for distinguishing classes
- **Confusion Matrix:** Breakdown of true/false positives/negatives

## Regression

(Metrics)

- $R^2$  Score: Proportion of variance explained by model
- MAE, RMSE: Average prediction errors

## Best Practice

Always report multiple metrics and visualize them for every model.



# 04

## Interpreting Model Metrics



# ● Why does metric choice matter?



Different applications value metrics differently

---

Recall matters in health Diagnosis

---

Precision in fraud detection

## ● How to flag issues?



### **High accuracy + low recall**

Model misses many positives (false negatives)



### **High recall + low precision**

Model creates many false alarms (false positives)

## ● What's a “good” score?



Depends on dataset, business context, and risk tolerance



05

## Monitoring Metrics During Training and Inference

# ● Training Phase Monitoring

01

Loss curves

Watch for convergence  
or divergence

02

Accuracy over  
epochs

Detect  
over/underfitting

03

Early stopping

Halt training if  
validation loss  
increases

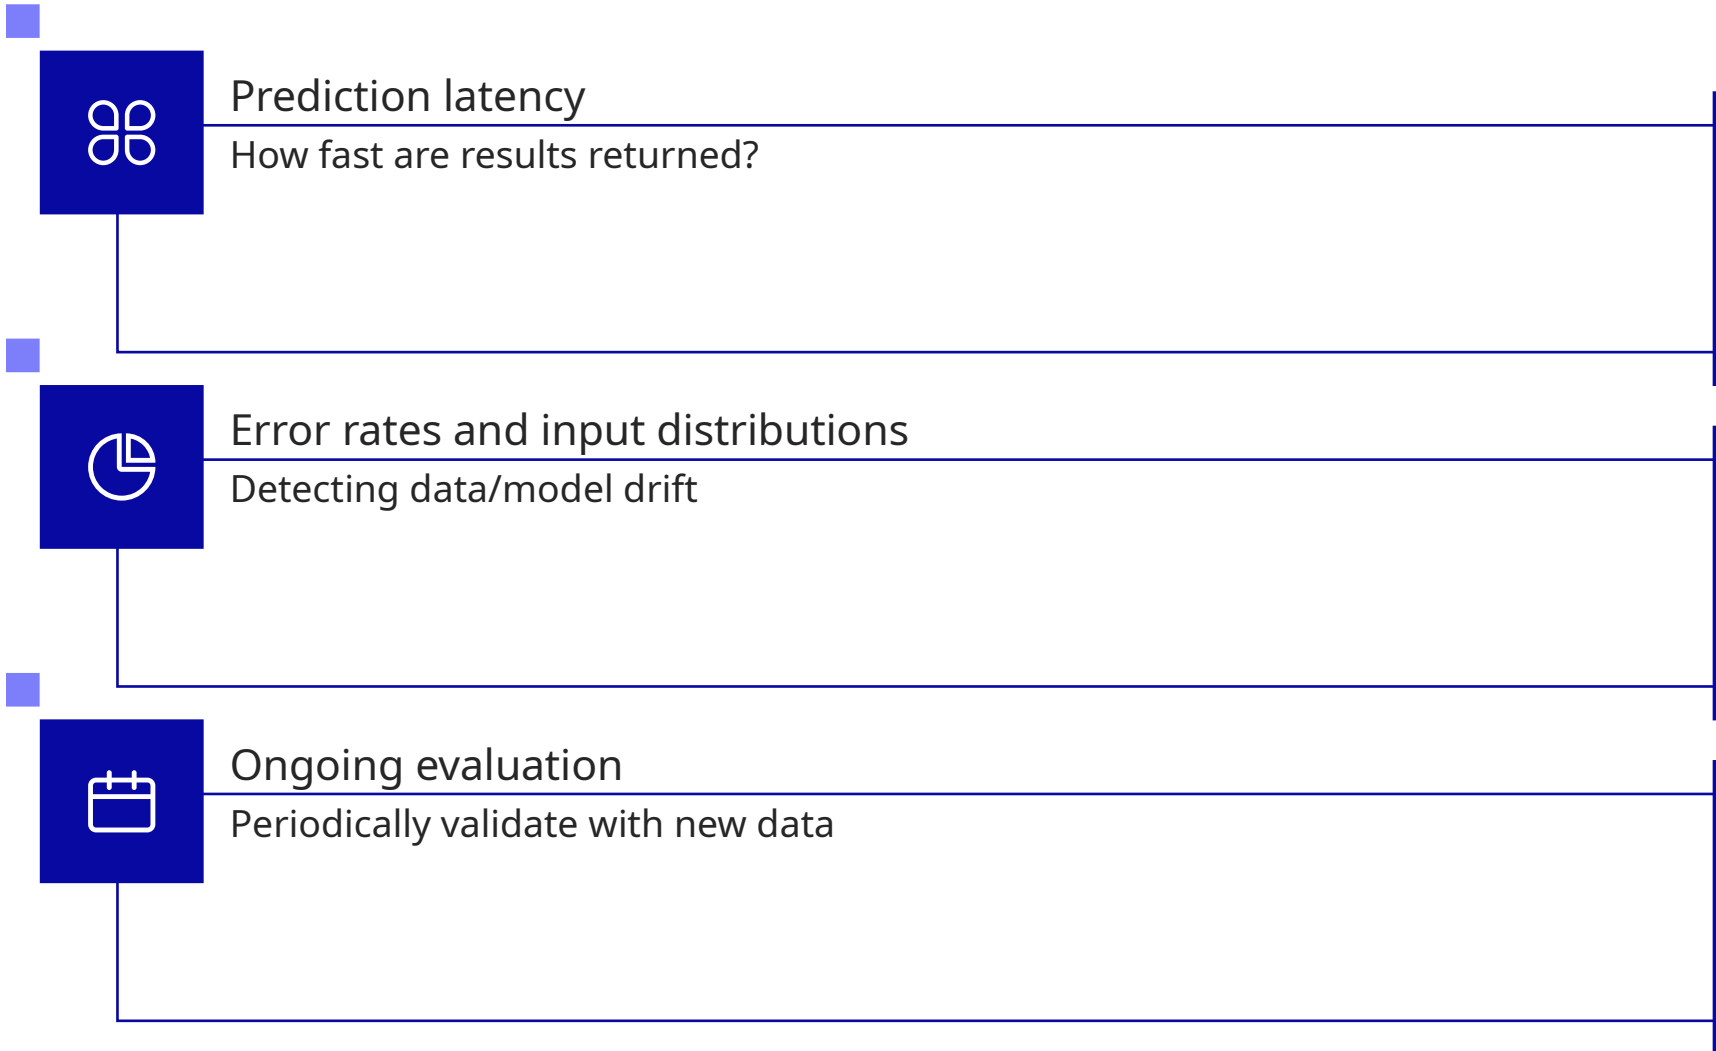
04

Time per  
epoch/batch

Performance  
monitoring



# ● Inference (Production) Monitoring





06

## Introduction to Model Deployment


# ● Goal



Make model available for real-world use

01

# ● Why deploy?



Enable predictions via APIs or user-facing applications  
Integrate into business processes

# Deployment Methods

—  
01

REST API (Flask/FastAPI)

—  
02

Web app (Streamlit/Gradio)

—  
03

Batch prediction






07

## Deploying with Flask – RESTAPI Approach

# ● Flask



01

Lightweight Python web framework

## ● How it works

”



Model loaded as a Python object



Receives HTTP requests with  
input data



Returns predictions as JSON

# ● Sample Flask API Structure

01

## Example Code

```
@app.route('/predict', methods=['POST'])
def predict():
    features = request.json['features']
    prediction = model.predict([features]).tolist()
    return jsonify({'prediction': prediction})
```



02

## Use Case

Easy integration with applications and microservices





08

## Deploying with Streamlit – Interactive Web App



# Deploying with Streamlit – Interactive Web App

## 01 Streamlit:

Tool for turning Python scripts into shareable web apps

## 02 Why Streamlit?

- Easily Display predictions, charts, explanations
- Instant web UIs with minimal code

## ● Sample Streamlit UI



### Example Code

```
import streamlit as st
input_val = st.number_input('Enter a value')
if st.button('Predict')pred = model.predict([[input_val]])
st.write(f"Prediction: {pred[0]}")
```

### Use Case:

Demos, prototyping, data science reporting



09

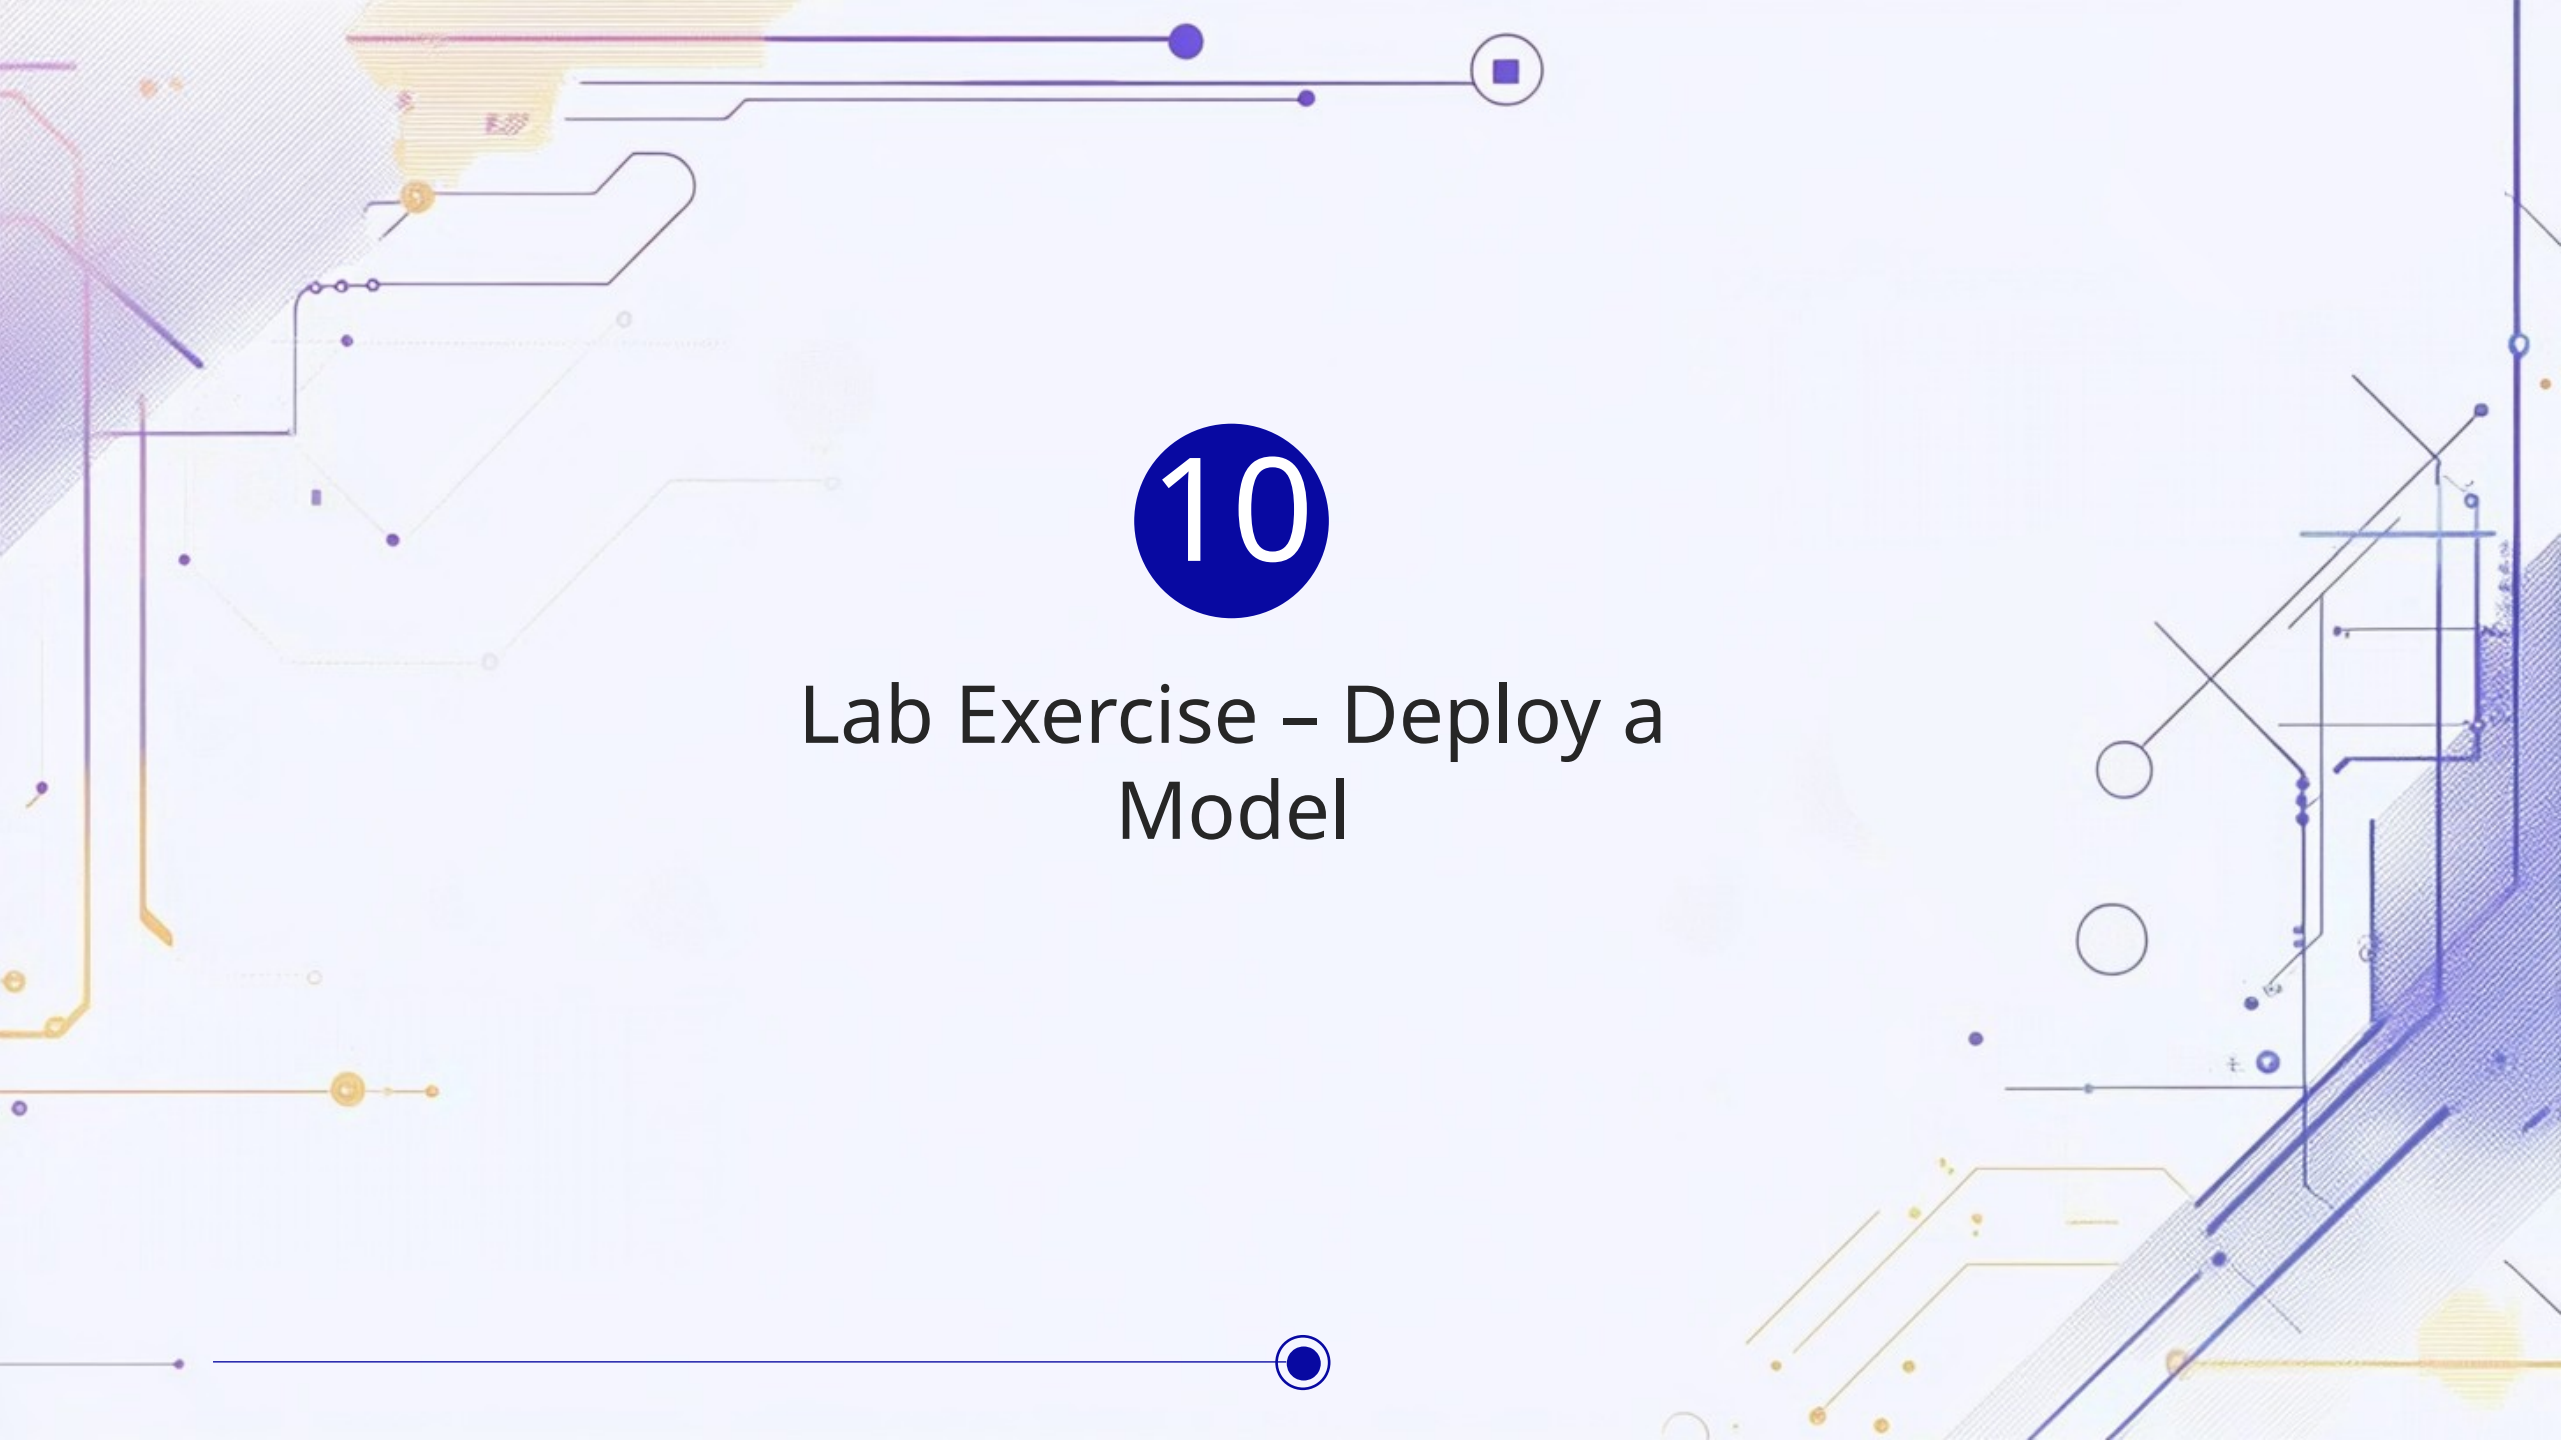
## Comparing Deployment Patterns

# ● Flask

01

|               | Flask                     | Streamlit                       |
|---------------|---------------------------|---------------------------------|
| Type          | REST API                  | Web App                         |
| Users         | Developers, other systems | End-user, business stakeholders |
| Interactivity | API calls only            | Full GUI, sliders, charts       |
| Best for...   | Production integration    | Demos, rapid prototyping        |





# 10


## Lab Exercise – Deploy a Model



# ● Objective



Make trained model accessible through Flask API and Streamlit app



# ● Lab Steps



Prepare your trained model  
Pickle/joblib save from  
Day 1 or Day 2



## Flask Deployment

Write a /predict  
endpoint  
Test with curl or  
Postman



## Streamlit Deployment

Create input form for  
user features  
Display prediction and  
basic metrics



## Test & Validate

Try several prediction  
examples  
Confirm outputs are  
correct and responsive



# Thanks

Karthikeyan Vaiyapuri